Sets of Half-Average Nulls Generate Risk-Limiting Audits: SHANGRLA

Philip B. Stark

University of California, Berkeley

Abstract. Risk-limiting audits (RLAs) for many social choice functions can be reduced to testing sets of null hypotheses of the form "the average of this list is not greater than 1/2" for a collection of finite lists of nonnegative numbers. Such social choice functions include majority, super-majority, plurality, multi-winner plurality, Instant Runoff Voting (IRV), Borda count, approval voting, and STAR-Voting, among others. The audit stops without a full hand count iff all the null hypotheses are rejected. The nulls can be tested in many ways. Ballot polling is particularly simple; two new ballot-polling riskmeasuring functions for sampling without replacement are given. Ballot-level comparison audits transform each null into an equivalent assertion that the mean of re-scaled tabulation errors is not greater than 1/2. In turn, that null can then be tested using the same statistical methods used for ballot polling—applied to different finite lists of nonnegative numbers: the SHANGRLA approach reduces auditing different social choice functions and different audit methods to the same simple statistical problem. Moreover, SHANGRLA comparison audits are more efficient than previous comparison audits for two reasons: (i) for most social choice functions, the conditions tested are both necessary and sufficient for the reported outcome to be correct, while previous methods tested conditions that were sufficient but not necessary, and (ii) the tests avoid a conservative approximation. The SHANGRLA abstraction simplifies stratified audits, including audits that combine ballot polling with ballot-level comparisons, producing sharper audits than the "SUITE" approach. SHANGRLA works with the "phantoms to evil zombies" strategy to treat missing ballot cards and missing or redacted cast vote records. That also facilitates sampling from "ballot-style manifests," which can dramatically improve efficiency when the audited contests do not appear on every ballot card. Open-source software implementing SHANGRLA ballot-level comparison audits is available. SHANGRLA was tested in a process pilot audit of an instant-runoff contest in San Francisco, CA, in November, 2019.

Keywords: sequential tests, martingales, Kolmogorov's inequality

Acknowledgments: I am grateful to Andrew Conway, Steven N. Evans, Kellie Ottoboni, Ronald L. Rivest, Vanessa Teague, Poorvi Vora, and Damjan Vukcevic for helpful conversations and comments on earlier drafts, and to Filip Zagorski for presenting the paper at the 5th Workshop on Advances in Secure Electronic Voting. The SHANGRLA software was a collaborative effort that included Michelle Blom, Andrew Conway, Dan King, Laurent Sandrolini, Peter Stuckey, and Vanessa Teague.

1 Introduction

A risk-limiting audit (RLA) of a reported election contest outcome is any procedure that guarantees a minimum probability of correcting the reported outcome if the reported

Philip B. Stark

2

winner(s) did not really win, but never alters a correct reported outcome. The largest probability that the procedure fail to correct the reported outcome if the reported outcome is wrong is the *risk limit*.

RLAs were introduced by [18] and named by [21]. RLA methods have been developed for a variety of social choice functions, to accommodate election equipment with different capabilities, and to comport with the logistics of ballot handling, organization, and storage in different jurisdictions.

RLAs are considered the gold standard for post-election tabulation audits, recommended by the National Academies of Science, Engineering, and Medicine [14], the Presidential Commission on Election Administration [17], the American Statistical Association [1], the League of Women Voters, Verified Voting Foundation, the Brennan Center for Justice, and other organizations concerned with election integrity.

Experience with RLAs. RLAs have been piloted dozens of times in 11 U.S. states and in Denmark. They are required by statute in Colorado, Nevada, Rhode Island, and Virginia, and authorized by statute in California and Washington.

Resources for RLAs. There is free and open-source software to help with the random selection of ballot cards for RLAs and to perform the risk calculations to determine when and if the audit can stop. ¹

Prerequisites for RLAs. A risk-limiting audit of a trustworthy paper record of voter intent can catch and correct wrong election outcomes. (Here, "trustworthy" means that a full hand count of the votes in the record would show the true winners.) But RLAs themselves check only the tabulation, not the trustworthiness of the paper records. If the paper trail is not trustworthy, it is not clear what a risk-limiting audit procedure accomplishes: while it can offer assurances that tabulation error did not alter the reported outcome, it cannot determine whether the reported outcome is right or wrong, nor can it promise any probability of correcting wrong outcomes. It therefore cannot limit the risk of certifying an outcome that is incorrect—the "risk" that a risk-limiting audit is supposed to limit.

Because all electronic systems are vulnerable to bugs, misconfiguration, and hacking, the paper trail can be trustworthy only if elections are conducted using *voter-verified paper ballots kept demonstrably secure* throughout the canvass and the audit. In particular, *compliance audits* [4,25,12,23] are needed to assure that the paper trail remains inviolate from the time of casting through the completion of the audit.

However, that is not enough: the means of marking the paper matters. Recent experiments have shown that voters rarely check BMD printout and rarely notice errors in BMD printout, even when instructed to check—verbally, through written instructions, and through signage [6]. Most BMD output is evidently not voter-verified. Moreover, while voters who notice errors in BMD printout may request a fresh chance to mark a ballot, there is

¹ See, e.g., https://www.stat.berkeley.edu/users/stark/Vote/auditTools.htm, https://www.stat.berkeley.edu/users/stark/Vote/ballotPollTools.htm, https://github.com/pbstark/auditTools, https://github.com/pbstark/CORLA18/blob/master/code/suite_toolkit.ipg and https://github.com/votingworks/arlo (all last visited 10 November 2019); an implementation of SHANGRLA ballot-level comparison audits is available at https://github.com/pbstark/SHANGRLA.

no way to tell whether the error was the voter's fault or the BMD's fault [10,5,2]. As a result, malfunctioning BMDs may go undetected by election officials. Applying RLA procedures to BMD printout cannot limit the risk that incorrect reported outcomes will go uncorrected, unless one simply defines "correct outcome" to be whatever an accurate tally of the paper would show, whether or not that reflects what voters indicated to the equipment.

What's new here. SHANGRLA uses a new abstract framing of RLAs that involves constructing a set of assertions for each contest. If the assertions are true, the contest outcomes are correct. The assertions are predicates on the set of ballot cards, that is, they are either true or false, depending on what votes the whole set of trusted paper ballot cards shows.

Each assertion is characterized by an assorter, a function that assigns a nonnegative number to each ballot card , again depending on the votes reflected on the ballot. The assertions that characterize the audit are of the form "the average value of the assorter for all the cast ballots is greater than 1/2." In turn, each assertion is checked by testing the complementary null hypothesis that the average is less than or equal to 1/2. To reject the entire set of complementary null hypotheses is to confirm the outcomes of all the contests under audit. Hence, the name of this method: Sets of Half-Average Nulls Generate RLAs (SHANGRLA).

By reducing auditing to repeated instances of a single, simple statistical problem—testing whether the mean of a list of nonnegative numbers is less than 1/2—SHANGRLA puts ballot-polling audits, comparison audits, batch-level comparison audits, and stratified and "hybrid" audits on the same footing, and puts auditing a broad range of social choice functions on the same footing. Moreover, it makes it easy to incorporate statistical advances into RLAs: only one function needs to be updated.

Open-source software implementing SHANGRLA audits is available.³ The software also implements the "phantoms to evil zombies" approach of [3] for dealing with missing cast-vote records and missing ballot cards. That also makes it possible to sample from "ballot-style manifests" [4,11], which facilitates efficient audits of contests that do not appear on every ballot card cast in the election. Despite the fact that they were developed more than 7 years ago, neither "phantoms-to-zombies" nor sampling from ballot-style manifests has been implemented before. The SHANGRLA software was tested in practice in a process pilot audit in San Francisco, CA, in November 2019.

2 Assorted Simplifications

An assorter A assigns a nonnegative value to each ballot card, depending on the marks the voter made on that ballot card.⁴

For instance, suppose that Alice and Bob are running against each other in a two-candidate first-past-the-post contest. The following function is an assorter: assign the

² A ballot consists of one or more ballot cards. Below, "ballot," "card," and "ballot card" are used interchangeably, even though in most U.S. jurisdictions, a ballot consists of more than one ballot card.

³ https://www.github.com/pbstark/SHANGRLA, last visited 22 November 2019.

⁴ The value might also depend on what the voting system reported for that ballot card and others. See section 3.2.

4 Philip B. Stark

value "1" to a ballot card if it has a mark for Alice but not for Bob; assign the value "0" if the card has a mark for Bob but not for Alice; assign the value 1/2, otherwise (e.g., if the card has an overvote or an undervote in this contest or does not contain the contest). Then Alice beat Bob iff the average value of the assorter for the full set of cast ballot cards is greater than 1/2: then Alice got more than 50% of the valid votes.

To express this more mathematically, let b_i denote the *i*th ballot card, and suppose there are N ballot cards in all. Let $1_{\text{Alice}}(b_i) = 1$ if ballot *i* has a mark for Alice, and 0 if not; define $1_{\text{Bob}}(b_i)$ analogously. The assorter could be written

$$A_{\text{Alice,Bob}}(b_i) \equiv \frac{1_{\text{Alice}}(b_i) - 1_{\text{Bob}}(b_i) + 1}{2}.$$

If b_i shows a mark for Alice but not for Bob, $A_{\text{Alice},\text{Bob}}(b_i) = 1$. If it shows a mark for Bob but not for Alice, $A_{\text{Alice},\text{Bob}}(b_i) = 0$. If it shows marks for both Alice and Bob (an overvote), for neither Alice nor Bob (an undervote), or if the ballot card does not contain the Alice v. Bob contest at all, $A_{\text{Alice},\text{Bob}}(b_i) = 1/2$. The average value of A over all ballot cards is

$$\bar{A}_{\mathrm{Alice,Bob}}^{b} \equiv \frac{1}{N} \sum_{i=1}^{N} A_{\mathrm{Alice,Bob}}(b_i).$$

If Alice is the reported winner, the contest can be audited at risk limit α by testing the complementary null hypothesis that $\bar{A}_{\rm Alice,Bob}^b \leq 1/2$ at significance level α . To reject the complementary null hypothesis is to conclude that Alice really won. If the complementary null hypothesis is true, Bob won or the contest was a tie: the assertion is false.

An assorter offers more flexibility than just counting votes. For instance, instead of either giving Alice one vote, half a vote, or no vote, an assorter could interpret a ballot card as giving Alice an arbitrary nonnegative value as a vote, depending on the marks on the ballot. This flexibility lets assorters solve the problem of auditing more complicated social choice functions, as we shall see.

2.1 Plurality elections

In a plurality contest with $K \ge 1$ winners and C > K candidates in all, a collection of candidates $\{w_k\}_{k=1}^K$ are the true winners and the remaining C - K candidates $\{\ell_j\}_{j=1}^{C-K}$ are the true losers iff the assertions

$$\bar{A}^b_{w_k,\ell_j} > 1/2, \quad \text{for all } 1 \leq k \leq K, \ \ 1 \leq j \leq C-K$$

all hold, where A_{w_k,ℓ_j} is defined as above, with w_k taking the role of Alice and ℓ_j taking the role of Bob. (This is essentially the approach taken in [18,20,22,12], reformulated and in different notation.) The contest can be audited to risk limit α by testing the K(C-K) hypotheses $\bar{A}^b_{w_k,\ell_j} \leq 1/2$ individually at significance level α . The audit stops only if all K(C-K) complementary hypotheses are rejected; otherwise it requires a full hand count.

2.2 Approval Voting

Even though the voting rules are different, the same assorter functions can be used to audit approval voting and plurality voting. Candidates $\{w_k\}_{k=1}^K$ are the winners and the

remaining C-K candidates $\{\ell_j\}_{j=1}^{C-K}$ are the losers of a K-winner approval voting contest iff all the assertions

$$\bar{A}_{w_k,\ell_j}^b > 1/2$$
, for all $1 \le k \le K$, $1 \le j \le C - K$

hold.

2.3 Super-majority

Suppose that a candidate must get at least a fraction $f \in (0,1)$ of the valid votes to win.⁵ [18] shows how to audit this social choice function, but it can also be expressed in terms of the assertion that the average of an assorter applied to the cast ballot cards is greater than 1/2.

Alice really won a super-majority contest with required winning fraction f iff

(valid votes for Alice) $> f \times$ ((valid votes for Alice) + (valid votes for everyone else)).

Define an assorter as follows:

$$A(b_i) \equiv \begin{cases} \frac{1}{2f}, & b_i \text{ has a mark for Alice and no one else} \\ 0, & b_i \text{ has a mark for exactly one candidate and not Alice} \\ 1/2, \text{ otherwise.} \end{cases}$$
 (1)

This assigns a nonnegative number to every ballot. Suppose that a fraction p > f of the valid votes are for Alice, and that a fraction q of the ballots have valid votes. Then

$$\bar{A}^b \equiv pq/(2f) + (1-q)/2 \ge q/2 + (1-q)/2 = 1/2.$$

Again, using assorters reduces auditing to the question of whether the average of a list of nonnegative numbers is greater than 1/2. The correctness of the outcome is implied by a single assertion, unlike plurality elections, which require (number of winners)×(number of losers) assertions.⁶

2.4 D'Hondt and other proportional representation schemes

[24] show how to reduce the problem of auditing D'Hondt and other proportional representation social choice functions to the problem of auditing a collection of two-candidate plurality contests. We have seen above that each such two-candidate contest can be expressed as the assertion that the average of an assorter applied to the ballots is greater than 1/2, so auditing proportional representation contests can be reduced to auditing a collection of assertions that the averages of a set of assorters over the cast ballots is greater than 1/2.

⁵ Values $f \le 1/2$ are not technically "super-majorities," but the generality is useful. For instance, the rules of some primaries in the U.S. eliminate candidates who receive less than 15% of the vote. An RLA using f = 0.15 might be used to check whether the correct candidates were eliminated.

⁶ To check whether K candidates all got at least a fraction $f \in (0,1)$ of the valid votes (with Kf < 1) requires testing at most K assertions.

2.5 Borda count, STAR-Voting, and other weighted additive voting schemes

Borda count is one of several voting systems that assign points to candidates for each ballot, depending on what the ballot shows; the winner is the candidate who receives the most points in total across all cast ballots. This involves only a slight generalization of plurality contests to account for the fact that a ballot can give a candidate more than one "point," while for plurality a candidate either gets a vote or not. As before, the reported result is correct if the reported winner actually received more points than each reported loser, which we can test by constructing an assorter for each (winner, loser) pair.

Let $s_{\text{Alice}}(b_i)$ denote a nonnegative "score" for Alice on ballot i, and let $s_{\text{Bob}}(b_i)$ be the score for Bob. These need not be integers. Let s^+ be an upper bound on the score any candidate can get on a ballot. Alice beat Bob iff

$$\sum_{i=1}^{N} s_{\text{Alice}}(b_i) > \sum_{i=1}^{N} s_{\text{Bob}}(b_i),$$

i.e., iff

$$\bar{s}_{\mathrm{Alice}}^b > \bar{s}_{\mathrm{Bob}}^b.$$

Make the affine transformation

$$A(b_i) \equiv \frac{s_{\text{Alice}}(b_i) - s_{\text{Bob}}(b_i) + s^+}{2s^+}.$$

Then $A(b_i) \geq 0$ and $\bar{s}_{Alice}^b > \bar{s}_{Bob}^b$ iff $\bar{A}^b > 1/2$.

2.6 Ranked-Choice and Instant-Runoff Voting (RCV/IRV)

[7] show how to reduce the correctness of a reported IRV winner to the correctness of the reported winners of a set of two-candidate plurality contests. The "candidates" in those contests are not necessarily the candidates in the original contest; they are just two mutually exclusive (but not exhaustive) conditions that a ballot might satisfy.

Two types of assertions can be combined to give sufficient conditions for the reported winner of an IRV contest to have really won:

- 1. Candidate i has more first-place ranks than candidate j has total mentions.
- 2. After a set of candidates have been eliminated from consideration, candidate i is the top ranked candidate on more ballot cards than candidate j is.

Both of these can be written as $\bar{A}^b>1/2$ by labeling the corresponding vote patterns "Alice" or "Bob" or "neither."

For instance, consider the first type of assertion. If b_i has candidate i ranked 1, the ballot is considered a vote for Alice. If b_i ranks candidate j at all, the ballot is considered a vote for Bob. Otherwise, the ballot is not a vote for either of them. If Alice beat Bob, candidate j cannot have beaten candidate i in the IRV contest.

In contrast to plurality, supermajority, approval, Borda, and d'Hondt, the assertions derived by [7] are sufficient for the reported winner to have won, but not necessary. Hence, it might be possible to sharpen such audits.

3 Auditing assertions

We audit the assertion $\bar{A}^b > 1/2$ by testing the complementary null hypothesis $\bar{A}^b \leq 1/2$ statistically. We audit until either all complementary null hypotheses about a contest are rejected at significance level α or until all ballots have been tabulated by hand. This yields a RLA of the contest in question at risk limit α .

3.1 Ballot-polling audits

Ballot-polling audits select individual ballot cards at random, either with or without replacement. The BRAVO method of [13] uses Wald's Sequential Probability Ratio Test (SPRT) for sampling with replacement.

For each (reported winner, reported loser) pair, BRAVO tests the conditional probability that a ballot contains a vote for the reported winner given that it contains a vote for the reported winner or the reported loser. Using assorters allows us to eliminate the conditioning and opens the door to a broader collection of statistical tests, including tests based on sampling without replacement, which can improve the efficiency of the audit. Two such methods are presented below. In contrast to the SPRT (with or without replacement), these methods only require knowing the reported winners, not the reported vote shares.

First, we shall derive the Kaplan-Kolmogorov method for sampling without replacement, based on ideas in Harold Kaplan's (now defunct) website. The method is based on the observation that a suitably constructed sequence is a nonnegative martingale, to which Kolmogorov's inequality for optionally stopped closed martingales can be applied.

We sample without replacement from a finite population of N nonnegative items, $\{x_1,\ldots,x_N\}$, with $x_j \geq 0$, $\forall j$. The population mean is $\mu \equiv \frac{1}{N} \sum_{j=1}^N x_j \geq 0$ and the population total is $N\mu \geq 0$. The value of the jth item drawn is X_j . On the hypothesis that $\mu = t$, $\mathbb{E}X_1 = t$, so $\mathbb{E}(X_1/t) = 1$. Conditional on X_1,\ldots,X_n , the total of the remaining N-n items is $N\mu - \sum_{j=1}^n X_j$, so the mean of the remaining items is

$$\frac{Nt - \sum_{j=1}^{n} X_j}{N - n} = \frac{t - \frac{1}{N} \sum_{j=1}^{n} X_j}{1 - n/N}.$$

Thus, the expected value of X_{n+1} given X_1, \ldots, X_n is $\frac{t-\frac{1}{N}\sum_{j=1}^n X_j}{1-n/N}$. Define

$$Y_1(t) \equiv \begin{cases} X_1/t, & Nt > 0, \\ 1, & Nt = 0, \end{cases}$$

and for $1 \le n \le N-1$,

$$Y_{n+1}(t) \equiv \begin{cases} X_{n+1} \cdot \frac{1 - \frac{n}{N}}{t - \frac{1}{N} \sum_{j=1}^{n} X_j}, & \sum_{j=1}^{n} X_j < Nt, \\ 1, & \sum_{j=1}^{n} X_j \ge Nt. \end{cases}$$

Then $\mathbb{E}(Y_{n+1}(t)|Y_1,\ldots Y_n)=1$. Let $Z_n(t)\equiv\prod_{j=1}^nY_j(t)$. Note that $Y_k(t)$ can be recovered from $\{Z_j(t),j\leq k\}$, since $Y_k(t)=Z_k(t)/Z_{k-1}(t)$. Now $\mathbb{E}|Z_k|\leq \max_j x_j<\infty$ and

$$\mathbb{E}(Z_{n+1}(t)|Z_1(t), \dots Z_n(t)) = \mathbb{E}(Y_{n+1}(t)Z_n(t)|Z_1(t), \dots Z_n(t)) = Z_n(t).$$

⁷ http://web.archive.org/web/20131209044835/http://printmacroj.com/martMean.htm

Thus

$$(Z_1(t), Z_2(t), \ldots, Z_N(t))$$

is a nonnegative closed martingale. By Kolmogorov's inequality, an application of Markov's inequality to martingales [8, p242] for any p > 0 and any $J \in \{1, ..., N\}$,

$$\Pr\left(\max_{1\leq j\leq J} Z_j(t) > 1/p\right) \leq p \, \mathbb{E}|Z_J|.$$

Since (Z_j) is a nonnegative martingale, $\mathbb{E}|Z_J| = \mathbb{E}Z_J = \mathbb{E}Z_1 = 1$. Thus a P-value for the hypothesis $\mu = t$ based on data $X_1, \dots X_J$ is $(\max_{1 \le j \le J} Z_j(t))^{-1} \wedge 1$.

However, if $X_j = 0$ for some j, then $Z_k = 0$ for all $k \geq j$. To avoid that problem, we can shift everything to the right: pick $\gamma > 0$, find a lower confidence bound for $\delta = \mu + \gamma > 0$ from data $\{X_j + \gamma\}$, then subtract γ from the lower confidence bound to get a lower confidence bound for μ . There are tradeoffs involved in picking γ : if many X_j turn out to be small, especially for small j, it helps to have γ large, and vice versa.

Unpacking the math yields the P-value

$$p_{KK} \equiv 1 \wedge \left(\max_{1 \le j \le J} \prod_{k=1}^{j} (X_k + \gamma) \frac{1 - (k-1)/N}{t - \frac{1}{N} \sum_{\ell=1}^{k-1} (X_\ell + \gamma)} \right)^{-1}$$

for the hypothesis that $\mu \leq t - \gamma$. This is implemented in the SHANGRLA software.

A related test that uses sampling without replacement, also introduced without proof on Kaplan's website, can be derived as follows. Let $S_j \equiv \sum_{k=1}^j X_k$, $\tilde{S}_j \equiv S_j/N$, and $\tilde{j} \equiv 1 - (j-1)/N$. Define

$$Y_n \equiv \int_0^1 \prod_{j=1}^n \left(\gamma \left[X_j \frac{\tilde{j}}{t - \tilde{S}_{j-1}} - 1 \right] + 1 \right) d\gamma.$$

This is a polynomial in γ of degree at most n, with constant term 1. Each X_j appears linearly. Under the null hypothesis that the population total is Nt, $\mathbb{E}X_1 = t$, and

$$\mathbb{E}(X_j \mid X_1, \dots, X_{j-1}) = \frac{Nt - S_{j-1}}{N - j + 1} = \frac{t - \tilde{S}_{j-1}}{\tilde{j}}.$$

Now

$$Y_1 = \int_0^1 \left(\gamma[X_1/t - 1] + 1 \right) d\gamma = \left[(\gamma^2/2)[X_1/t - 1] + \gamma \right]_{\gamma=0}^1 = [X_1/t - 1]/2 + 1 = \frac{X_1}{2t} + 1/2.$$

Thus, under the null,

$$\mathbb{E}Y_1 = \frac{\mathbb{E}X_1}{2t} + 1/2 = 1.$$

Also,

$$\mathbb{E}(Y_n|X_1,\dots,X_{n-1}) = \mathbb{E}\left[\int_0^1 \prod_{j=1}^n \left(\gamma \left[X_j \frac{\tilde{j}}{t - \tilde{S}_{j-1}} - 1\right] + 1\right) d\gamma \middle| X_1,\dots,X_{n-1}\right]$$

$$= \int_0^1 \left(\gamma \left[\mathbb{E}(X_n|X_1,\dots,X_{n-1}) \frac{\tilde{n}}{t - \tilde{S}_{n-1}} - 1\right] + 1\right) \prod_{j=1}^{n-1} \left(\gamma \left[X_j \frac{\tilde{j}}{t - \tilde{S}_{j-1}} - 1\right] + 1\right) d\gamma$$

$$= \int_0^1 \left(\gamma \left[\frac{t - \tilde{S}_{n-1}}{\tilde{n}} \frac{\tilde{n}}{t - \tilde{S}_{n-1}} - 1\right] + 1\right) \prod_{j=1}^{n-1} \left(\gamma \left[X_j \frac{\tilde{j}}{t - \tilde{S}_{j-1}} - 1\right] + 1\right) d\gamma$$

$$= \int_0^1 \prod_{i=1}^{n-1} \left(\gamma \left[X_j \frac{\tilde{j}}{t - \tilde{S}_{j-1}} - 1 \right] + 1 \right) d\gamma = Y_{n-1}.$$

Hence, under the null hypothesis, $(Y_j)_{j=1}^N$ is a nonnegative closed martingale with expected value 1, and Kolmogorov's inequality implies that for any $J \in \{1, ..., N\}$,

$$\Pr\left(\max_{1\leq j\leq J} Y_j(t) > 1/p\right) \leq p.$$

This method for finding a P-value for the hypothesis $\bar{A}^b \leq 1/2$ is also implemented in the SHANGRLA software, using a novel approach to integrating the polynomial recursively due to Steven N. Evans (U.C. Berkeley). This was the risk-measuring function used in the San Francisco IRV process pilot audit in November, 2019.

3.2 Ballot-comparison audits

Ballot-comparison audits require the voting system to export a cast vote record (CVR) for each physical ballot—the system's interpretation of that ballot—in such a way that the corresponding physical ballot can be retrieved, and vice versa. Suppose that we apply the assorters for a contest to the CVRs (rather than the actual physical ballots), and every assorter has an average greater than 1/2, i.e., the assertions are true for the CVRs. Then the assertions are true for the physical ballots provided the CVRs did not inflate the average value of the assorter by more than the assorter margin, twice the mean of the assorter applied to the reported CVRs, minus 1, as we shall see.

By how much could error in an individual CVR inflate the value of the assorter compared to the value the assorter has for the actual ballot? Since the assorter does not assign a negative value to any ballot, the *overstatement error* ω_i for CVR i is at most the value the assorter assigned to CVR i.

The existence of an upper bound for the overstatement error is key to auditing: otherwise, a single extreme value could make a contest result wrong, and it would take a prohibitively large sample to rule out that possibility.

If we can reject the hypothesis that the mean overstatement error for an assorter is large enough to account for the assorter margin, we may conclude that the assorter that the assorter mean exceeds 1/2 is true, and the audit of that assertion can stop.

Let b_i denote the *i*th ballot, and let c_i denote the cast-vote record for the *i*th ballot. Let A denote an assorter, which maps votes on a ballot card or on a CVR into [0, u], where u is an upper bound on the value A assigns to any ballot card or CVR. For instance, for plurality voting, u = 1; for super-majority, u = 1/(2f), where f is the fraction of valid votes required to win.

The $overstatement\ error$ for the ith ballot is

$$\omega_i \equiv A(c_i) - A(b_i) \le A(c_i) \le u. \tag{2}$$

It is the amount by which the assorter applied to the cast vote record overstates the value of the assorter applied to corresponding physical ballot.

⁸ However, see section sec. 3.4 below.

Let

$$\bar{A}^c \equiv \frac{1}{N} \sum_{i=1}^N A(c_i) \text{ and } \bar{\omega} \equiv \frac{1}{N} \sum_{i=1}^N \omega_i.$$

Now $\bar{A}^b = \bar{A}^c - \bar{\omega}$, so $\bar{A}^b > 1/2$ iff $\bar{\omega} < \bar{A}^c - 1/2$. We know that $\bar{A}^c > 1/2$ (or the assertion would not be true for the CVRs), so $\bar{\omega} < \bar{A}^c - 1/2$ iff

$$\frac{\bar{\omega}}{2\bar{A}^c - 1} < 1/2.$$

Define $v \equiv 2\bar{A}^c - 1$, the reported assorter margin. In a two-candidate plurality contest, v is the fraction of ballot cards with valid votes for the reported winner, minus the fraction with valid votes for the reported loser. This is the diluted margin of [22,12]. (Margins are traditionally calculated as the difference in votes divided by the number of valid votes. Diluted refers to the fact that the denominator is the number of ballot cards, which is greater than or equal to the number of valid votes.)

With this notation, the condition for the assertion to be true is:

$$\frac{\bar{\omega}}{v} < 1/2.$$

Let $\tau_i \equiv 1 - \omega_i/u \ge 0$, and $\bar{\tau} \equiv (1/N) \sum_{i=1}^N \tau_i = 1 - \bar{\omega}/u$. Then $\bar{\omega} = u(1 - \bar{\tau})$ and

$$\frac{\bar{\omega}}{v} = \frac{u}{v}(1 - \bar{\tau}).$$

Now $\omega/v < 1/2$ iff $\frac{u}{v}(1-\bar{\tau}) < 1/2$, i.e.,

$$-\frac{u}{v}\bar{\tau} < 1/2 - \frac{u}{v}$$

$$\bar{\tau} > 1 - \frac{v}{2u}$$

$$\frac{\bar{\tau}}{2 - \frac{v}{u}} > 1/2.$$

Finally, define

$$B(b_i, c) \equiv \tau_i/(2 - v/u) = \frac{1 - \omega_i/u}{2 - v/u} > 0, \ i = 1, \dots, N.$$

Then B assigns nonnegative numbers to ballots, and the outcome is correct iff

$$\bar{B} \equiv \frac{1}{N} \sum_{i=1}^{N} B_i > 1/2.$$

It is an assorter! Any technique that can be used with ballot polling, including those in sec. 3.1, can also be used to test the assertion $\bar{B} > 1/2$.

This assertion-based approach to comparison audits is sharper than methods that rely on the maximum across-contest relative overstatement of margins (MACRO) [22] in at least two ways: avoiding combining overstatements across candidates or contests gives a sharper upper bound on the total error in each (winner, loser) sub-contest, and the test statistic avoids combining overstatements across contests, which otherwise can lead to unnecessary escalation of the audit if an overstatement is observed in a contest with a wide margin.

3.3 Stratified audits

Stratified sampling can be useful within a jurisdiction if the jurisdiction has a heterogenous mix of election equipment with different capability, e.g., ballot-polling for precinct-based optical scan, where it can be difficult to associate ballot cards with cast-vote records, and ballot-level comparisons for central-count optical scan. Stratification can also be useful auditing contests that cross jurisdictional boundaries by allowing those jurisdictions to sample independently from each other.

Stratified sampling for RLAs has been addressed in a number of papers, including [18,19,20,9,16]. The central idea of the approach taken in SUITE [16] can be used with SHANGRLA to accommodate stratified sampling and to combine ballot-polling and ballot-level comparison audits: Look at all allocations of error across strata that would result in an incorrect outcome. Reject the hypothesis that the outcome is incorrect if the maximum P-value across all such allocations is less than the risk limit.

SHANGRLA will generally yield a sharper (i.e., more efficient) test than SUITE, because it deals more efficiently with ballot cards that do not contain the contest in question, because it avoids combining overstatements across candidate pairs and across contests, and because it accommodates sampling without replacement more efficiently.

With SHANGRLA, whatever the sampling scheme used to select ballots or groups of ballots, the underlying statistical question is the same: is the average value of each assorter applied to all the ballot cards greater than 1/2?

Suppose the cast ballot cards are partitioned into $S \geq 2$ strata, where stratum s contains N_s ballot cards, so $N = \sum_{s=1}^{S} N_s$. Let \bar{A}_s^b denote the mean of the assorter applied to just the ballot cards in stratum s. Then

$$\bar{A}^b = \frac{1}{N} \sum_{s=1}^{S} N_s \bar{A}^b_s = \sum_{s=1}^{S} \bar{A}^b_s \frac{N_s}{N}.$$

We can reject the hypothesis $\bar{A}^b \leq 1/2$ if we can reject the hypothesis

$$\bigcap_{s \in S} \left\{ \bar{A}_s^b \frac{N_s}{N} \le \beta_s \right\}$$

for all $(\beta_s)_{s=1}^S$ such that $\sum_{s=1}^S \beta_s \leq 1/2$. Let $P_s(\beta_s)$ be a P-value for the hypothesis $\bar{A}_s^b \leq \frac{N}{N_s} \beta_s$. That P-value could result from ballot polling, ballot-level comparison, batch-level comparison, or any other valid method. For instance, it could be produced by the methods described in sec. 3.1 or sec. 3.2.

Suppose that the samples from different strata are independent, so that $\{P_s(\beta_s)\}_{s=1}^S$ are independent random variables. Then Fisher's combining function (or any other method for nonparametric combination of tests) can be used to construct an overall P-value for the intersection null hypothesis $\bigcap_{s\in S} \left\{\frac{N_s}{N}\bar{A}_s^b \leq \beta_s\right\}$. In particular, if the intersection hypothesis is true, then the probability distribution of

$$-2\sum_{s=1}^{S}\ln P_s(\beta_s)$$

is dominated by the chi-square distribution with 2S degrees of freedom, as discussed in [16]. That makes it possible to assign a conservative P-value to the hypothesis $\bigcap_{s \in S} \left\{ \frac{N_s}{N} \bar{A}_s^b \leq \beta_s \right\}$ for every $(\beta_s)_{s=1}^S$ such that $\sum_{s=1}^S \beta_s \leq 1/2$. If all such S-tuples can be rejected, we may conclude that $\bar{A}_b > 1/2$.

3.4 Zombie Bounds II: Return of the Missing Ballot

[3] discuss how to conduct RLAs when not every ballot is accounted for or when a ballot cannot be retrieved. They cover both ballot-polling audits and ballot-level comparison audits. This section presents a brief but more systematic treatment for ballot-level comparison audits, reflecting what the SHANGRLA software implements. This method also makes it possible to use ballot-card style information to target the sampling to ballot cards that the voting system claims contain the contest, while protecting against the possibility that the voting system does not report that information accurately.

To conduct a RLA, it is crucial to have an upper bound on the total number of ballot cards cast in the contest. Absent such a bound, arbitrarily many ballots could be missing from the tabulation, and the true winner(s) could be any candidate(s). Let N denote an upper bound on the number of ballot cards that contain the contest. Suppose that $n \leq N$ CVRs contain the contest and that each of those CVRs is associated with a unique, identifiable physical ballot card that can be retrieved if that CVR is selected for audit. The phantoms-to-evil zombies approach is as follows.

If N>n, create N-n "phantom ballots" and N-n "phantom CVRs." Calculate the assorter mean for all the CVRs—including the phantoms—treating the phantom CVRs as if they contain no valid vote in the contest contest (i.e., the assorter assigns the value 1/2 to phantom CVRs). Find the corresponding assorter margin (twice the assorter mean, minus 1).

To conduct the audit, sample integers between 1 and N.

- If the resulting integer is between 1 and n, retrieve and inspect the ballot card associated with the corresponding CVR.
 - If the associated ballot contains the contest, calculate the overstatement error as in equation {eq. 2}.
 - If the associated ballot does not contain the contest, calculate the overstatement error using the value the assorter assigned to the CVR, but as if the value the assorter assigns to the physical ballot is zero (that is, the overstatement error is equal to the value the assorter assigned to the CVR).
- If the resulting integer is between n+1 and N, we have drawn a phantom CVR and a phantom ballot. Calculate the overstatement error as if the value the assorter assigned to the phantom ballot was 0 (turning the phantom into an "evil zombie"), and as if the value the assorter assigned to the CVR was 1/2.

Proposition: if the risk is calculated based on this substitution of "evil zombies" for "phantoms," the result is still a RLA with risk limit α .

Proof: Every unaccounted for ballot card that might have or should have contained the contest is treated in the least favorable way. Every unaccounted for CVR is treated in exactly the way it was tabulated by the assorter, namely, it is assigned the value 1/2.

Some jurisdictions, notably Colorado, redact CVRs if revealing them might compromise vote anonymity. If such CVRs are omitted from the tally and the number of phantom CVRs and ballots are increased correspondingly, this approach still leads to a valid RLA. But if they are included in the tally, then if they are selected for audit they should be treated as if they had the value u (the largest value the assorter can assign) in calculating the overstatement error.

4 Discussion

4.1 From many, one

Even though SHANGRLA may involve testing many assertions in the audit of one or more contests, there is no need to adjust for *multiplicity*. If any assertion is false, the chance that its complementary hypothesis will be rejected is at most α . If more than one assertion is false, the chance that all the complementary hypotheses will be rejected is at most α , because the probability of the intersection of a collection of events cannot exceed the probability of any event in the collection. Thus, if any of the reported winners did not really win, the chance that every complementary null hypothesis will be rejected is at most α : the chance that the audit will stop without a full hand count is not greater than the risk limit.

4.2 Sharpness and Efficiency

Extant comparison audit methods rely on MACRO, the maximum across-contest relative overstatement of margins [19]. MACRO is embedded in Colorado's CORLA audit tool, in the Arlo audit tool, and https://www.stat.berkeley.edu/~stark/Vote/auditTools.htm. MACRO involves combining discrepancies across pairs of candidates and across contests in a way that is conservative, but not sharp. That is, the condition that is tested is necessary for one or more reported outcomes to be incorrect, but not sufficient. In contrast, by keeping the pairwise margins separate, SHANGRLA is sharp for plurality, super-majority, approval, Borda, STAR-Voting, D'Hondt, etc.—but in general not for RCV/IRV. The conditions it tests are both necessary and sufficient for one or more outcomes to be incorrect. This generally allows smaller sample sizes to confirm the results when the reported contest outcomes are correct.

4.3 The Power of Positivity

Working with assertions reduces election auditing to testing hypotheses of the form $\bar{A}^b < 1/2$: the only statistical issue is to test whether the mean of a finite list of nonnegative numbers is less than 1/2. As new techniques for testing that hypothesis are developed, they can be applied immediately to election audits.

4.4 To Halve or Halve not?

Assertions might look more elegant expressed as $\bar{A}^b > 1$ rather than $\bar{A}^b > 1/2$, which would just involve re-defining A by multiplying it by 2. However, keeping the connection between assorter means and getting more than 50% of the vote in a two-candidate majority contest seems to be a helpful mnemonic. It also might feel more natural to write an assertion as or $\bar{A}^b > 0$, but that would cut the connection to getting more than a 50% vote share and make the lower bound less natural than the nonnegativity constraint $A(b_i) > 0$ for all i.

Similarly, defining the "assorter margin" to be $v \equiv 2(\bar{A}^c - 1/2)$ rather than $\bar{A}^c - 1/2$ keeps the parallel to a two-candidate plurality contest, where the "margin" would generally be defined to be the winner's vote share minus the loser's vote share.

5 Conclusions

Risk-limiting audits of a broad variety of social choice functions can be reduced to testing whether the mean of any list among a set of finite lists of nonnegative numbers is less than or equal to 1/2. That is, Sets of Half-Average Nulls Generate Risk-Limiting Audits (SHANGRLA). Those hypotheses can be tested directly, e.g., by ballot polling, or indirectly, by ballot-level comparisons or other methods. They can also be tested using Bernoulli sampling [15], stratified sampling, and "hybrid" methods following the same general approach as SUITE [16] (see 3.3), but SHANGRLA is generally more efficient. The sampling unit can be an individual ballot, or a cluster of ballots (e.g., all ballots cast in a single precinct or tabulated by a particular machine). Samples can be drawn with or without replacement.

Ballot-level comparison audits can also be framed as testing whether any of the means of a set of finite lists of nonnegative numbers is less than or equal to 1/2, allowing exactly the same statistical tests to be used for ballot-polling audits and for ballot-level comparison audits, unifying the treatment of audits.

This paper proves the validity of two hypothesis tests for that statistical problem based on sampling without replacement, both of which were stated without proof in a now-defunct website of Harold Kaplan but apparently never published. Both proofs are based on Kolmogorov's inequality for optionally stopped martingales.

Even though auditing one or more contests generally involves testing many half-average nulls, no multiplicity adjustment is needed, because the audit only stops if all the nulls are rejected.

For many social choice functions (including plurality, multi-winner plurality, majority, super-majority, approval, D'Hondt, Borda count, and STAR-Voting), SHANGRLA comparison audits are sharper than previous comparison audit methods based on MACRO because the conditions it tests are both necessary and sufficient for the reported outcomes to be wrong, while previous methods tested conditions that were necessary but not sufficient. (MACRO bounds the maximum of a set of sums by the sum of the term-by-term maxima, both in the condition and in the test statistic; SHANGRLA keeps the maxima separate, both in the condition and in the test statistic.)

SHANGRLA also "plays nice" with the phantoms-to-zombies approach [3] for dealing with missing ballot cards and missing cast-vote records, which has two benefits: (i) it makes it easy to treat missing ballots rigorously, and (ii) it can substantially improve the efficiency of auditing contests that do not appear on every ballot card, by allowing the sample to be drawn just from cards that the voting system claims contain the contest, without having to trust that the voting system correctly identified which cards contain the contest.

Open-source software implementing SHANGRLA comparison audits and the phantoms-to-zombies approach is available; the software was tested in a process pilot audit of an IRV contest in San Francisco, CA, in November 2019.

References

1. American Statistical Association: American Statistical Association statement on risk-limiting post-election audits.

- www.amstat.org/outreach/pdfs/Risk-Limiting_Endorsement.pdf (2010)
- Appel, A., DeMillo, R., Stark, P.: Ballot-marking devices cannot assure the will of the voters. https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3375755 (2019)
- Bañuelos, J., Stark, P.: Limiting risk by turning manifest phantoms into evil zombies. Tech. rep., arXiv.org (2012), http://arxiv.org/abs/1207.3413, retrieved 17 July 2012
- 4. Benaloh, J., Jones, D., Lazarus, Ε., Lindeman, Μ., Stark, Secrecy-preserving observableSOBA: ballot-level audits. In: Proceed-2011 Electronic Voting Technology Workshop Elections (EVT/WOTE Trustworthy '11). USENIX http://statistics.berkeley.edu/~stark/Preprints/soba11.pdf
- Bernhard, M., Halderman, J., Rivest, R., Vora, P., Ryan, P., Teague, V., Benaloh, J., Stark, P., Wallach, D.: Public evidence from secret ballots. In: Krimmer, R., Volkamer, M., Binder, N.B., Kersting, N., Pereira, O., Schürmann, C. (eds.) Electronic Voting. E-Vote-ID 2017. Lecture Notes in Computer Science, 10615. Springer (2017), https://doi.org/10.1007/978-3-319-68687-5_6
- 6. Bernhard, M., McDonald, A., Meng, H., Hwa, J., Bajaj, N., Chang, K., Halderman, J.: Can voters detect malicious manipulation of ballot marking devices? In: 41st IEEE Symposium on Security and Privacy. p. (in press). IEEE (2020)
- 7. Blom, M., Stuckey, P.J., Teague, V.: Risk-limiting audits for irv elections. arXiv preprint arXiv:1903.08804 (2019), https://arxiv.org/abs/1903.08804
- 8. Feller, W.: An introduction to probability theory and its applications. Vol. II. John Wiley & Sons, Inc., New York (1971)
- Higgins, M., Rivest, R., Stark, P.: Sharper p-values for stratified post-election audits. Statistics, Politics, and Policy 2(1) (2011), http://www.bepress.com/spp/vol2/iss1/7
- Kaczmarek, T., Wittrock, J., Carback, R., Florescu, A., Rubio, J., Runyan, N., Vora, P., Zagorski, F.: Dispute resolution in accessible voting systems: The design and use of audiotegrity. In: Heather, J., Schneider, S., Teague, V. (eds.) E-Voting and Identity: 4th International Conference, Vote-ID 2013, Guildford, UK, July 2013. LNCS Volume 7985, pp. 127–141. Springer (2013)
- 11. Lindeman, M., McBurnett, N., Ottoboni, K., Stark, P.: Next steps for the Colorado risk-limiting audit (CORLA) program. https://arxiv.org/abs/1803.00698 (2018)
- 12. Lindeman, M., Stark, P.: A gentle introduction to risk-limiting audits. IEEE Security and Privacy ${f 10},\ 42–49\ (2012)$
- 13. Lindeman, M., Stark, P., Yates, V.: BRAVO: Ballot-polling risk-limiting audits to verify outcomes. In: Proceedings of the 2011 Electronic Voting Technology Workshop / Workshop on Trustworthy Elections (EVT/WOTE '11). USENIX (2012)
- 14. National Academies of Sciences, Engineering, and Medicine: Securing the Vote: Protecting American Democracy. The National Academies Press, Washington, DC (Sep 2018). https://doi.org/10.17226/25120
- 15. Ottoboni, K., Bernhard, M., Halderman, A., Rivest, R., Stark, P.: Bernoulli ballot polling: A manifest improvement for risk-limiting audits. Voting '19 (2019), preprint:

http://arxiv.org/abs/1812.06361

- 16. Ottoboni, K., Stark, P., Lindeman, M., McBurnett, N.: Risk-limiting audits by stratified union-intersection tests of elections (suite). In: Electronic Voting. E-Vote-ID 2018. Lecture Notes in Computer Science. Springer (2018), https://link.springer.com/chapter/10.1007/978-3-030-00419-4_12, preprint: https://arxiv.org/abs/1809.04235
- The 17. Presidential Commission on Election Administration: American Voting Experience: Report and recommendations the Presidential Commission on Election Administration. https://www.eac.gov/assets/1/6/Amer-Voting-Exper-final-draft-01-09-14-508.pdf (2014)
- 18. Stark, P.: Conservative statistical post-election audits. Ann. Appl. Stat. 2, 550–581 (2008), http://arxiv.org/abs/0807.4005
- 19. Stark, P.: A sharper discrepancy measure for post-election audits. Ann. Appl. Stat. 2, 982–985 (2008), http://arxiv.org/abs/0811.1697
- Stark, P.: CAST: Canvass audits by sampling and testing. IEEE Transactions on Information Forensics and Security, Special Issue on Electronic Voting 4, 708–717 (2009)
- 21. Stark, P.: Risk-limiting post-election audits: P-values from common probability inequalities. IEEE Transactions on Information Forensics and Security 4, 1005–1014 (2009)
- 22. Stark, P.: Super-simple simultaneous single-ballot risk-limiting audits. In: Proceedings of the 2010 Electronic Voting Technology Workshop / Workshop on Trustworthy Elections (EVT/WOTE '10). USENIX (2010), http://www.usenix.org/events/evtwote10/tech/full_papers/Stark.pdf
- 23. Stark, P.: An introduction to risk-limiting audits and evidence-based elections, written testimony prepared for the Little Hoover Commission. https://www.stat.berkeley.edu/~stark/Preprints/lhc18.pdf (2018)
- 24. Stark, P.B., Teague, V.: Verifiable european elections: Risk-limiting audits for d'hondt and its relatives. JETS: USENIX Journal of Election Technology and Systems 3.1 (2014), https://www.usenix.org/jets/issues/0301/stark
- Stark, P.B., Wagner, D.A.: Evidence-based elections. IEEE Security and Privacy 10, 33–41 (2012)